

## MLP



Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 128)	100480
=====		
dense_1 (Dense)	(None, 128)	16512
=====		
dense_2 (Dense)	(None, 10)	1290
=====		
activation (Activation)	(None, 10)	0
=====		
Total params: 118,282		
Trainable params: 118,282		
Non-trainable params: 0		
=====		



*I used a two layer MLP with softmax activation and no regularization methods. It uses the Keras backend and can achieve just over 80% accuracy.*

Why isn't MLE a good method for loss here?

*The output space is discrete, so a notion of "distance to the correct answer" is nonsensical.*



*As one may observe from the preceding graphs, the accuracy does not seem to full converge and the model may benefit from additional epochs. In the future, this will be an improvement that I will consider.*





*These don't really resemble specific articles of clothing, but you can clearly see some symmetries in the weightings.*

## CNN

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 64)	640
-----		
activation (Activation)	(None, 28, 28, 64)	0
-----		
max_pooling2d (MaxPooling2D)	(None, 14, 14, 64)	0
-----		
flatten (Flatten)	(None, 12544)	0
-----		
dense (Dense)	(None, 128)	1605760
-----		
activation_1 (Activation)	(None, 128)	0
-----		
dense_1 (Dense)	(None, 10)	1290
-----		
activation_2 (Activation)	(None, 10)	0
=====		
Total params: 1,607,690		
Trainable params: 1,607,690		
Non-trainable params: 0		

*As you can see from the output above, my model uses a single convolution and pooling, followed by a flattening layer and two dense layers, with each step having an activation between.*



*As you can see from the above graphs, the model may benefit from additional epochs. I have therefore increased the number of epochs to 50 for the grading mode.*

How many matrices?

1

What are the dimensions?


28x28

After maxpooling?

27x27



## PCA



*I chose a component value of 350, because it looks like that explains most of the variance (the graph above should suffice for showing the error at each level, since explained variance and error are directly related). As you can see from the two graphs below, the PCA seems to help with convergence slightly (the second is with PCA):*

